

Introduction to Operating Systems


What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - ◆ Execute user programs and make solving user problems easier
 - ◆ Make the computer system convenient to use
 - ◆ Use the computer hardware in an efficient manner

Computer System Structure

- Computer system can be divided into four components:
 - ◆ **Hardware** – provides basic computing resources
 - CPU, memory, I/O devices
 - ◆ **Operating system**
 - Controls and coordinates use of hardware among various applications and users
 - ◆ **Application programs** – define the ways in which the system resources are used to solve the computing problems of the users
 - Word processors, compilers, web browsers, database systems, video games
 - ◆ **Users**
 - People, machines, other computers

Four Components of a Computer System



What Operating Systems Do?

- Users want convenience, **ease of use** and **good performance**
 - ◆ Don't care about **resource utilization**
- But shared computer such as **mainframe** or **minicomputer** must keep all users happy
- Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- Handheld computers are resource poor, optimized for usability and battery life
- Some computers have little or no user interface, such as embedded computers in devices and automobiles could have real-time requirements

So What are the Major OS Functions?

- **Control access** and provide interfaces
 - ◆ To the OS and devices attached to the system
 - ◆ Provide interfaces for human-machine and machine-machine transactions
- Manage resources (**Resource Allocator**)
 - ◆ Mediate resource usage among different tasks
 - ◆ Implement policies

So What are the Major OS Functions?


■ Provide abstractions

- ◆ Hide the peculiarities of the hardware.
- ◆ Example: device independent I/O

■ Consume resources*

- ◆ OS runs on the system that is being managed... so it is going to consume resources!



Hardware



Operating System needs to give an easy-to-use abstraction over a complicated collection of hardware components

Files

How does OS manage persistent storage (e.g., disks)?




OSes use the File concept to give a familiar way of managing the persistent storage

Storage Management: In Words

- OS provides uniform, logical view of information storage
 - ◆ Abstracts physical properties to logical storage unit - **file**
 - ◆ Each medium is controlled by device (i.e., disk drive, tape drive)
 - Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- File-System management
 - ◆ Files usually organized into directories
 - ◆ Access control on most systems to determine who can access what
 - ◆ OS activities include
 - Creating and deleting files and directories
 - Primitives to manipulate files and directories
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media

Storage Management: In Pictures




File system based
view – content visible
as files and directories

Block based view
content considered as
blocks of data

Physical view –
all data just a stream
of bits

Storage-Device Hierarchy



Storage Hierarchy



- Storage systems organized in hierarchy
 - ◆ Speed
 - ◆ Cost
 - ◆ Volatility
- **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage
- **Device Driver** for each device controller to manage I/O
 - ◆ Provides uniform interface between controller and kernel

Storage Structure

- Main memory – only large storage media that the CPU can access directly
 - ◆ **Random access**
 - ◆ Typically **volatile**
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
- Hard disks – rigid metal or glass platters covered with magnetic recording material
 - ◆ Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - ◆ The **disk controller** determines the logical interaction between the device and the computer
- **Solid-state disks** – faster than hard disks, nonvolatile
 - ◆ Various technologies
 - ◆ Becoming more popular

Programs

How to load a program?



Programs are made up of Code & Data
Need to load both into memory and make the Code run

Process Management


- A process is a program in execution. It is a unit of work within the system. Program is a *passive entity*, process is an *active entity*.
- Process needs resources to accomplish its task
 - ◆ CPU, memory, I/O, files
 - ◆ Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one program counter specifying location of next instruction to execute
 - ◆ Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs
 - ◆ Concurrency by multiplexing the CPUs among the processes / threads

Process Management Activities

The operating system is responsible for the following activities in connection with process management:


- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling

Concurrency




With multiple processes, we can think of each process getting a virtual processor


Parallelism




Virtual Processor




Virtual Processor




Virtual Processor



Processor



Processor

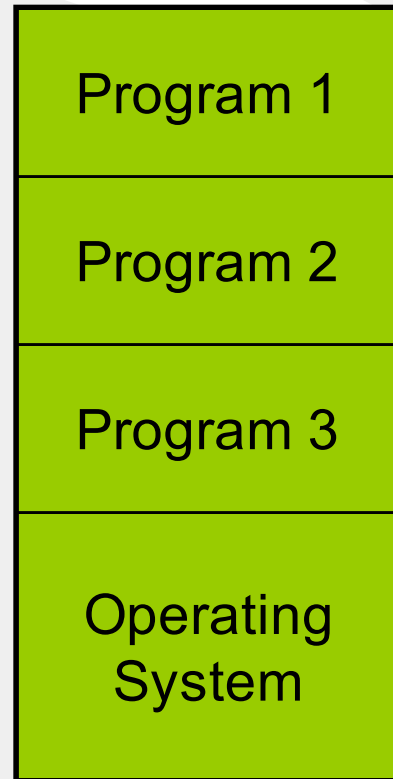


Processor



Parallelism

Memory Sharing



Memory

OS needs to share the memory among all the processes that are active so that a process does not corrupt memory belonging to another process

Memory Management

- To execute a program all (or part) of the instructions must be in memory
- All (or part) of the data that is needed by the program must be in memory.
- Memory management determines what is in memory and when
 - ◆ Optimizing CPU utilization and computer response to users
- Memory management activities
 - ◆ Keeping track of which parts of memory are currently being used and by whom
 - ◆ Deciding which processes (or parts thereof) and data to move into and out of memory
 - ◆ Allocating and deallocating memory space as needed

Operating System Definition


- No universally accepted definition
 - “Everything a vendor ships when you order an operating system” is a good approximation
 - ◆ But varies wildly
 - “The one program running at all times on the computer” is the **kernel**.
- Everything else is either
 - ◆ a system program (ships with the operating system) , or
 - ◆ an application program.

Computer Startup

- **bootstrap program** is loaded at power-up or reboot
 - ◆ Typically stored in ROM or EPROM, generally known as **firmware**
 - ◆ Initializes all aspects of system
 - ◆ Loads operating system kernel and starts execution

Computer System Org.

- Computer-system operation
 - ◆ One or more CPUs, device controllers connect through common bus providing access to shared memory
 - ◆ Concurrent execution of CPUs and devices competing for memory cycles



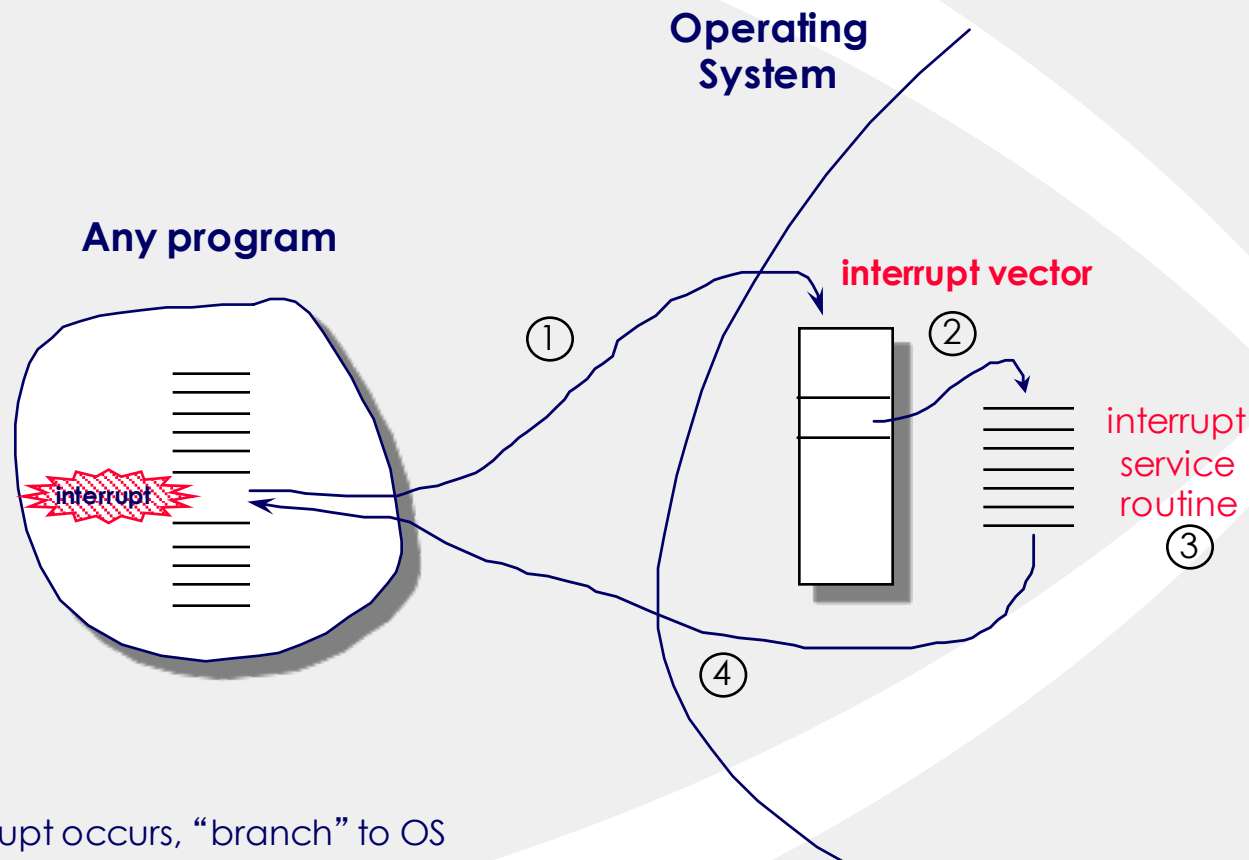
Computer-System Operation

- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an **interrupt**

Common Functions of Interrupts

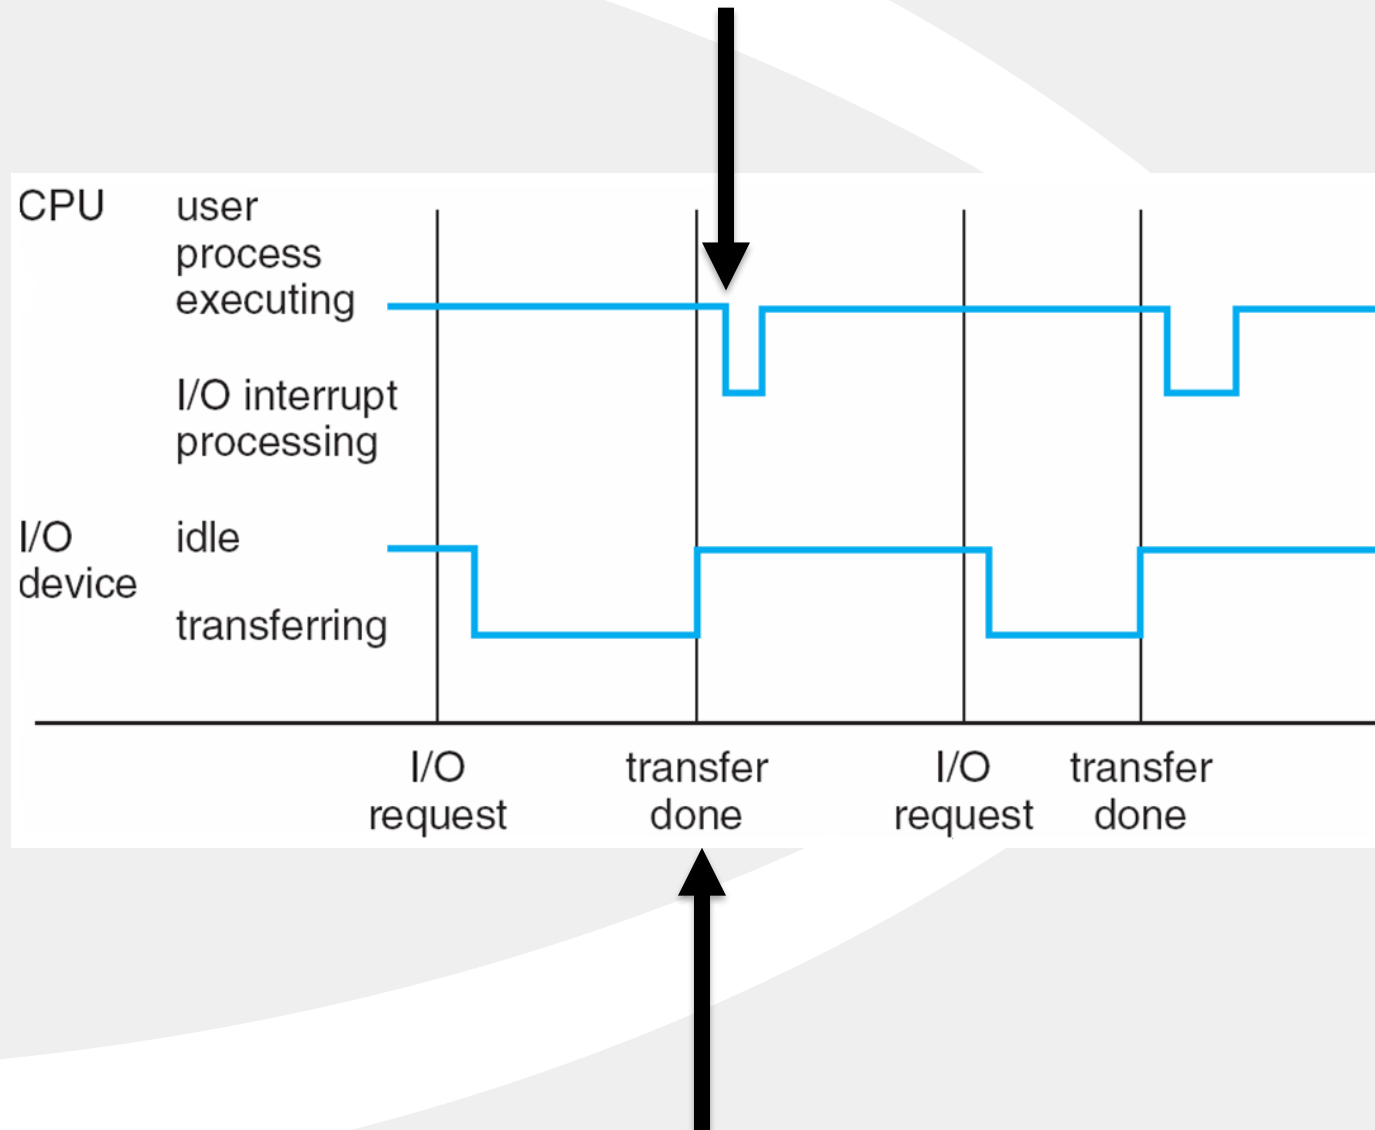
- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- An operating system is **interrupt driven**

Interrupt Servicing: In Picture




1. An interrupt occurs, "branch" to OS
2. Locate the interrupt service routine (ISR) via interrupt vector
3. Execute the ISR
4. Return to interrupted program

Interrupt Timeline



How a Modern Computer Works




A von Neumann architecture

Computer-System Arch.


- Started migrating from single processors to multiprocessors (multi-cores are common)
- **Special purpose processors** are also heavily used (e.g., GPUs)
 - ◆ Also known as **parallel systems, tightly-coupled systems**
 - ◆ Advantages include:
 1. **Increased throughput**
 2. **Economy of scale**
 3. **Increased reliability** – graceful degradation or fault tolerance
 - ◆ Two types:
 1. **Asymmetric Multiprocessing** – each processor is assigned a specific task.
 2. **Symmetric Multiprocessing** – each processor performs all tasks

Symmetric Multiprocessing Architecture



A Dual-Core Design


- Multi-chip and **multicore**
- Systems containing all chips
 - ◆ Chassis containing multiple separate systems



Clustered Systems

- Like multiprocessor systems, but multiple systems working together
 - ◆ Usually sharing storage via a **storage-area network (SAN)**
 - ◆ Provides a **high-availability** service which survives failures
 - **Asymmetric clustering** has one machine in hot-standby mode
 - **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - ◆ Some clusters are for **high-performance computing (HPC)**
 - Applications must be written to use **parallelization**
 - ◆ Some have **distributed lock manager (DLM)** to avoid conflicting operations

Clustered Systems



Operating System Structure

- **Multiprogramming** (**Batch system**) needed for efficiency
 - ❖ Single user cannot keep CPU and I/O devices busy at all times
 - ❖ Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - ❖ A subset of total jobs in system is kept in memory
 - ❖ One job selected and run via **job scheduling**
 - ❖ When it has to wait (for I/O for example), OS switches to another job

Operating System Structure

- **Timesharing (multitasking)** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - ◆ **Response time** should be < 1 second
 - ◆ Each user has at least one program executing in memory \Rightarrow **process**
 - ◆ If several jobs ready to run at the same time \Rightarrow **CPU scheduling**
 - ◆ If processes don't fit in memory, **swapping** moves them in and out to run
 - ◆ **Virtual memory** allows execution of processes not completely in memory

OS Operations

■ **Interrupt driven** (hardware and software)

- ◆ Hardware interrupt by one of the devices
- ◆ Software interrupt (**exception** or **trap**):
 - Software error (e.g., division by zero)
 - Request for operating system service
 - Other process problems include infinite loop, processes modifying each other or the operating system

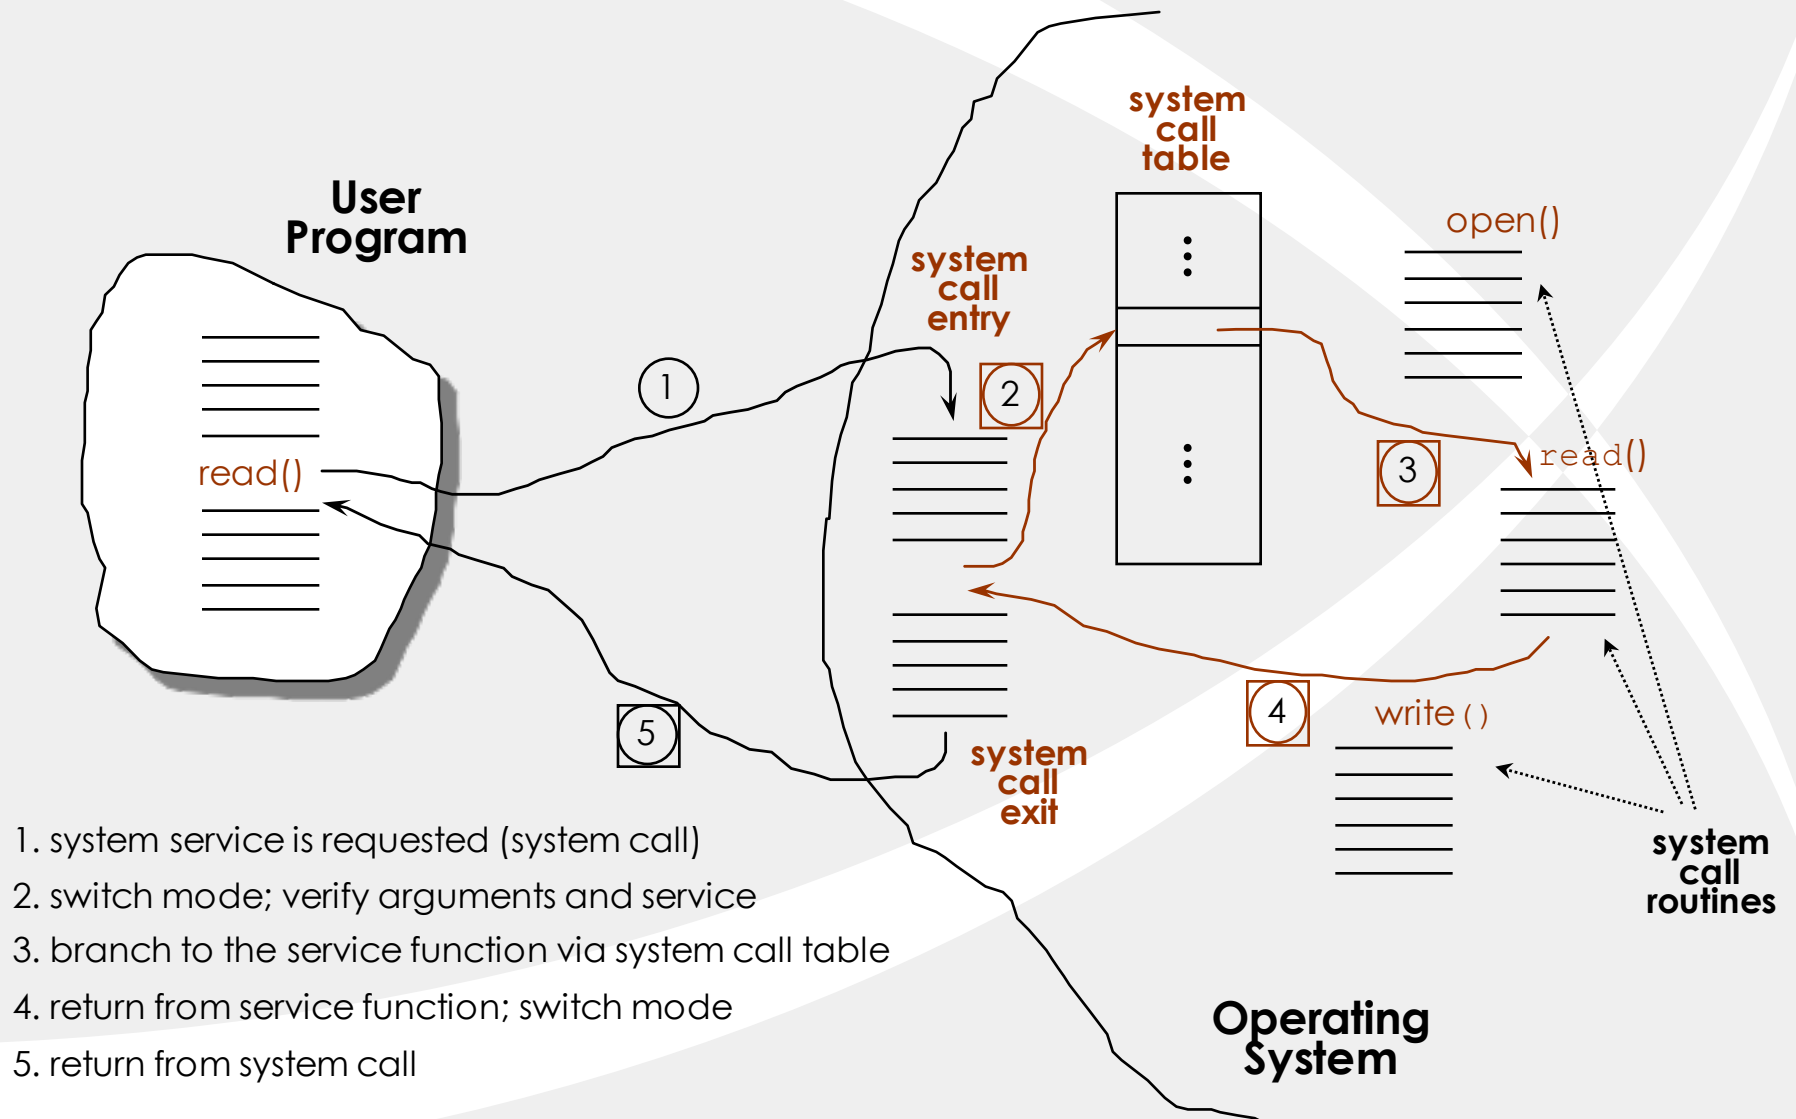
OS Operations (cont.)

- **Dual-mode** operation allows OS to protect itself and other system components
 - ◆ **User mode** and **kernel mode**
 - ◆ **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as **privileged**, only executable in kernel mode
 - System call changes mode to kernel, return from call resets it to user
- Increasingly CPUs support multi-mode operations
 - ◆ i.e. **virtual machine manager (VMM)** mode for guest **VMs**

User to Kernel Mode Transition

- Timer to prevent infinite loop / process hogging resources
 - ◆ Timer is set to interrupt the computer after some time period
 - ◆ Keep a counter that is decremented by the physical clock.
 - ◆ Operating system set the counter (privileged instruction)
 - ◆ When counter zero generate an interrupt
 - ◆ Set up before scheduling process to regain control or terminate program that exceeds allotted time

How system call is processed?



Protection and Security

- **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- **Security** – defense of the system against internal and external attacks
 - ◆ Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users, to determine who can do what
 - ◆ User identities (**user IDs**, security IDs) include name and associated number, one per user
 - ◆ User ID then associated with all files, processes of that user to determine access control
 - ◆ Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
 - ◆ **Privilege escalation** allows user to change to effective ID with more rights

Computing Environments - Traditional

- Stand-alone general purpose machines
- But blurred as most systems interconnect with others (i.e., the Internet)
- **Portals** provide web access to internal systems
- **Network computers** (**thin clients**) are like Web terminals
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks

Computing Environments – Distributed

- Distributed computing
 - ◆ Collection of separate, possibly heterogeneous, systems networked together
 - **Network** is a communications path, **TCP/IP** most common
 - **Local Area Network (LAN)**
 - **Wide Area Network (WAN)**
 - **Metropolitan Area Network (MAN)**
 - **Personal Area Network (PAN)**
 - ◆ **Network Operating System** provides features between systems across network
 - Communication scheme allows systems to exchange messages
 - Illusion of a single system

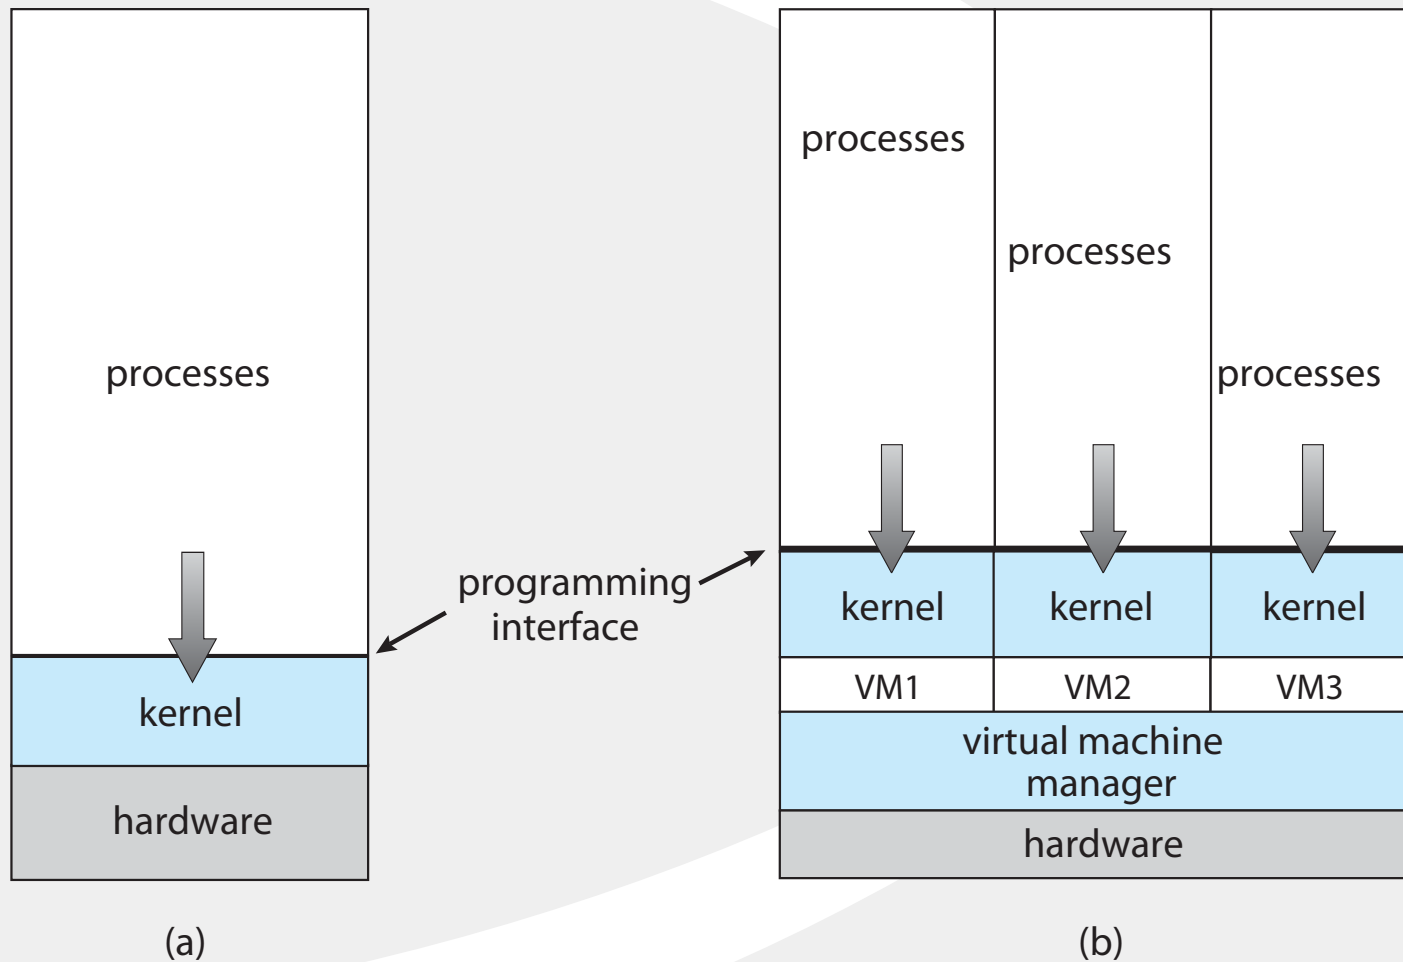
Computing Environments - Virtualization

- Allows operating systems to run applications within other OSes
 - ◆ Vast and growing industry
- **Emulation** used when source CPU type different from target type (i.e. PowerPC to Intel x86)
 - ◆ Generally slowest method
 - ◆ When computer language not compiled to native code – **Interpretation**
- **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - ◆ Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS
 - ◆ **VMM** (virtual machine Manager) provides virtualization services

Computing Environments - Virtualization

- Use cases involve laptops and desktops running multiple OSES for exploration or compatibility
 - ◆ Apple laptop running Mac OS X host, Windows as a guest
 - ◆ Developing apps for multiple OSES without having multiple systems
 - ◆ QA testing applications without having multiple systems
 - ◆ Executing and managing compute environments within data centers
- VMM can run natively, in which case they are also the host
 - ◆ There is no general purpose host then (VMware ESX and Citrix XenServer)

Computing Environments - Virtualization




Computing Environments – Cloud Computing

- Delivers computing, storage, even apps as a service across a network
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - ◆ Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage
- Many types
 - ◆ **Public cloud** – available via Internet to anyone willing to pay
 - ◆ **Private cloud** – run by a company for the company's own use
 - ◆ **Hybrid cloud** – includes both public and private cloud components
 - ◆ Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
 - ◆ Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - ◆ Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)

Computing Environments – Cloud Computing

- Cloud computing environments composed of traditional OSes, plus VMMs, plus cloud management tools
 - ◆ Internet connectivity requires security like firewalls
 - ◆ Load balancers spread traffic across multiple applications



Computing Environments – Real-Time Embedded Systems

- Real-time embedded systems most prevalent form of computers
 - ◆ Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - ◆ Use expanding
- Many other special computing environments as well
 - ◆ Some have OSes, some perform tasks without an OS
- Real-time OS has well-defined fixed time constraints
 - ◆ Processing **must** be done within constraint
 - ◆ Correct operation only if constraints met