

What is a *deadlock*?

- **Deadlock** -- *permanent* blocking of a set of processes that compete for system resources
- **No efficient solution** to the deadlock problem in the general case.
- In deadlock
 - ◆ a set of **processes are in a wait** state, because each process is waiting for a resource that is held by some other waiting process
- **All deadlocks involve conflicting resource needs by two or more processes**

Classification of resources—I

- Resources grouped into two classes:
 - ◆ **Reusable**: something that can be safely used by one process at a time and is ***not depleted*** by that use. Processes obtain resources that they later release for reuse by others.

E.g., CPU, memory, specific I/O devices, or files.
 - ◆ **Consumable**: these can be ***created and destroyed***. When a resource is acquired by a process, the resource ceases to exist.

E.g., interrupts, signals, or messages.

Classification of resources—II

- Another classification of resources:
 - ◆ **Preemptable**: these can be taken away from the process owning it with no ill effects (needs save/restore).
E.g., memory or CPU.
 - ◆ **Non-preemptable**: cannot be taken away from its current owner without causing the computation to fail.
E.g., printer or floppy disk.
- Deadlocks mostly occur when sharing **reusable** and **non-preemptable** resources.

Conditions for deadlock

- Four conditions that must hold for a deadlock to be possible:
 - ◆ **Mutual exclusion:** processes require exclusive control of its resources (not sharing)
 - ◆ **Hold and wait:** process may wait for a resource while holding others
 - ◆ **No preemption:** process will not give up a resource until it is finished with it
 - ◆ **Circular wait:** each process in the chain holds a resource requested by another

Conditions for deadlock

- ***Also, a process cannot be reset to an earlier state where resources not held***
- Conditions 1—3 are necessary, but not sufficient. All 4 are needed

Solving deadlocks

- If a necessary condition is prevented a deadlock need not occur. For example:
 - ◆ Systems with only shared resources cannot deadlock.
 - Negates *mutual exclusion*.
 - ◆ Systems that abort processes which request a resource that is in use.
 - Negates *hold and wait*.

Solving deadlocks (contd.)

- If a necessary condition is prevented a deadlock need not occur. For example:
 - ◆ Preemptions may be possible if a process does not use its resources until it has acquired all it needs.
 - Negates *no preemption*.
 - ◆ Systems that detect or avoid deadlocks.
 - Prevents *cycle*.
 - ◆ Transaction processing systems provide checkpoints so that processes may back out of a transaction.
 - Negates *irreversible process*.

Resource allocation graphs

Set of Processes



$$P = \{P_1, P_2, \dots, P_n\}$$

Set of Resources


$$R = \{R_1, R_2, \dots, R_m\}$$

Some resources come in multiple units.

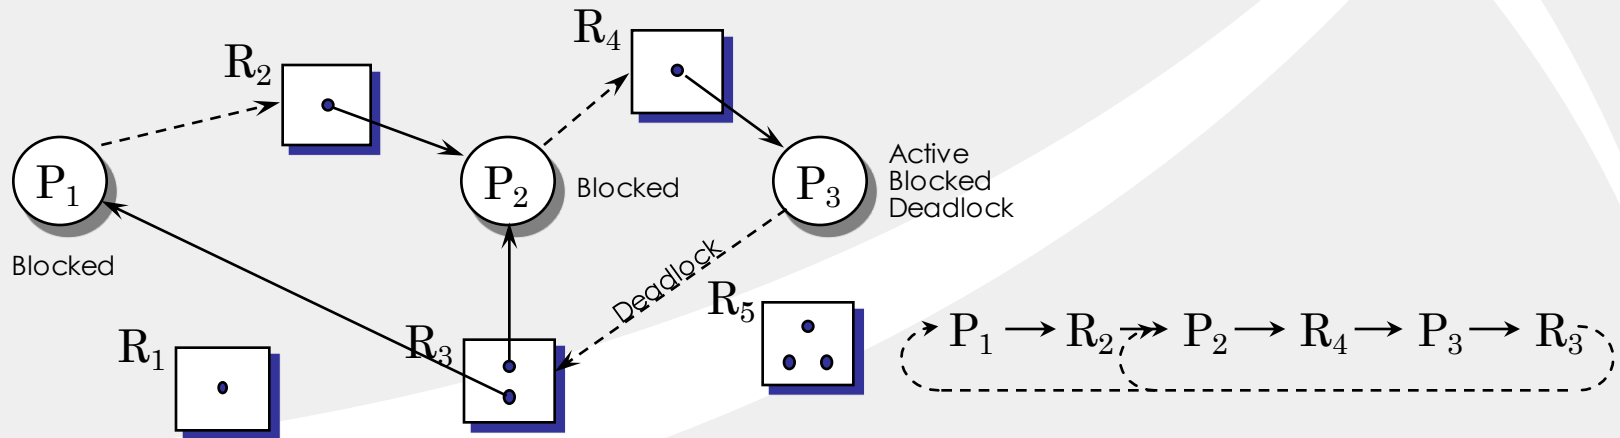
R_j has 2 units




Process P_i waits for (has requested) R_j



Resource R_j has been allocated to P_i



Cycle is necessary, but ...



Multiple resource unit case:


No Deadlock—yet!

Because, either P₂ or P₄ could relinquish a resource allowing P₁ or P₃ (which are currently blocked) to continue. P₂ is still executing, even if P₄ requests R₁.

... a knot is required

- Cycle is a *necessary condition* for a deadlock. With multiple unit resources—*not sufficient*.
- A *knot* must exist—a cycle with no non-cycle outgoing path from any involved node.
- At the moment assume that:
 - ◆ a process *halts* as soon as it waits for one resource, and
 - ◆ processes can wait for only *one* resource at a time

Further requests



(a) If P_4 requests R_1 no deadlock since P_2 may release R_1 , allowing P_4 to complete.

(b) $P_1 \rightarrow R_2 \rightarrow P_3 \rightarrow R_2 \rightarrow P_4 \rightarrow R_1 \rightarrow P_2 \rightarrow R_2$

(b) If P_2 requests R_2 : Deadlock—*Cycle—Knot*.

No active processes to release resources.

Strategies for deadlocks

- In general, four strategies are used for dealing with deadlocks:
 - ◆ **Ignorance**: pretend there is no problem at all.
 - ◆ **Prevention**: design a system in such a way deadlock is excluded a priori.
 - ◆ **Avoidance**: make a decision dynamically checking whether a request will, if granted, potentially lead to a deadlock or not.
 - ◆ **Detection**: let the deadlock occur and detect when it happens, and take some action to recover after the fact.

Dealing with Deadlocks

- Different people react to this strategy in different ways:
 - ◆ **Mathematicians:** find deadlock totally unacceptable, and say that it must be prevented at all costs.
 - ◆ **Engineers:** ask how serious it is, and do not want to pay a penalty in performance and convenience.
- UNIX approach
 - ◆ ignore the problem
 - ◆ Prevention price is high – inconvenient restrictions

Deadlock prevention

- Deadlock prevention is to design a system to exclude the possibility of a deadlock
 - ◆ *indirect methods* -- prevent the occurrence of one of the necessary conditions listed earlier.
 - ◆ *direct methods* -- prevent the occurrence of a circular wait condition.
- Deadlock prevention strategies are very **conservative** -- solve the problem of deadlock by limiting access to resources and by imposing restrictions on processes

More on deadlock prevention

◆ Mutual exclusion

- In general, this condition cannot be disallowed.

◆ Hold-and-wait

- The hold and-wait condition can be prevented by requiring that a process request all its required resources at one time, and blocking the process until all requests can be granted simultaneously.

◆ No preemption

- One solution is that if a process holding certain resources is denied a further request, that process must release its unused resources and request them again, together with the additional resource.

◆ Circular Wait

- The circular wait condition can be prevented by defining a linear ordering of resource types. If a process has been allocated resources of type R, then it may subsequently request only those resources of types following R in the ordering.

Deadlock avoidance

- Allows the necessary conditions
- Makes judicious resource allocation choices to ensure a deadlock-free system
 - ◆ System evaluates current request and denies it if granting it will lead to potential deadlock
 - ◆ Requires knowledge of future requests
- Deadlock avoidance approaches:
 - ◆ Resource trajectories
 - ◆ Safe/unsafe states
 - ◆ Dijkstra's Banker's algorithm

Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare the ***maximum number*** of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes


Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
- That is:
 - ◆ If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - ◆ When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - ◆ When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Basic Facts

- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

Safe, Unsafe, Deadlock State




Avoidance Algorithms

- Single instance of a resource type
 - ◆ Use a resource-allocation graph
- Multiple instances of a resource type
 - ◆ Use the banker's algorithm


Resource-Allocation Graph Scheme

- **Claim edge** $P_i \rightarrow R_j$ indicated that process P_i may request resource R_j ; represented by a dashed line
- Claim edge converts to request edge when a process requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system

Resource-Allocation Graph



Unsafe State In Resource-Allocation Graph



Resource-Allocation Graph Algorithm

- Suppose that process P_i requests a resource R_j
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

Banker's Algorithm

- Multiple instances
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

Banker's algorithm—*definitions*

Assume N Processes $\{P_i\}$
 M Resources $\{R_j\}$

Availability vector **Avail_j**, units of each resource (initialized to maximum, changes dynamically).

Let [**Max_{ij}**] be an $N \times M$ matrix.

Max_{ij} = L means Process P_i will request at most L units of R_j .

[**Hold_{ij}**] Units of R_j currently held by P_i

[**Need_{ij}**] Remaining need by P_i for units of R_j

Need_{ij} = **Max_{ij}** - **Hold_{ij}**, for all i & j

Banker's algorithm—*resource request*

- At any instance, P_i posts its request for resources in vector REQ_j .
- **Step 1:** verify that a process matches its needs.
■ *if* $REQ_j > Need_{ij}$ *abort—error, impossible*
- **Step 2:** check if the requested amount is available.
■ *if* $REQ_j > Avail_j$ *goto Step 1— P_i must wait*
- **Step 3:** provisional allocation.
■ $Avail_j = Avail_j - REQ_j$
■ $Hold_{ij} = Hold_{ij} + REQ_j$
■ $Need_{ij} = Need_{ij} - REQ_j$
■ *if* $isSafe()$ *then grant resources—system is safe*
■ *else cancel allocation; goto Step 1— P_i must wait*

Banker's algorithm—*isSafe*

Find out whether the system is in a safe state.

Work and **Finish** are two temporary vectors.

Step 1: initialize.

$\text{Work}_j = \text{Avail}_j$ for all j ; $\text{Finish}_i = \text{false}$ for all i .

Step 2: find a process P_i such that

$\text{Finish}_i = \text{false}$ and $\text{Need}_{ij} \leq \text{Work}_j$
if no such process, *goto Step 4*.

Step 3: $\text{Work}_j = \text{Work}_j + \text{Hold}_{ij}$

$\text{Finish}_i = \text{true}$

goto Step 2.

Step 4: *if* $\text{Finish}_i = \text{true}$ for all i

then return true—yes, the system is safe

else return false—no, the system is NOT safe

Banker's algorithm—*what is safe?*

■ Safe with respect to some resource allocation.

◆ very safe

$NEED_i \leq AVAIL$ for all Processes P_i .

Processes can run to completion in any order.

◆ safe (but take care)

$NEED_i > AVAIL$ for some P_i

$NEED_i \leq AVAIL$ for at least one P_i such that

There is at least one correct order in which the processes may complete their use of resources.

◆ unsafe (deadlock inevitable)

$NEED_i > AVAIL$ for some P_i

$NEED_i \leq AVAIL$ for at least one P_i

But some processes cannot complete successfully.

◆ deadlock

$NEED_i > AVAIL$ for all P_i

Processes are already blocked or will become so as they request a resource.

Example—safe allocation

	Max	Hold	Need	Finish	Avail	Work
P_1	5	2	3	F	2	2
P_2	4	1	3	F		
P_3	2	1	1	F		

For simplicity, assume that all the resources are identical.

Assume P_1 acquires one unit. *Very safe?* **No!** $\text{Need}_2 > 2$

Safe? Let us see with the safe/unsafe algorithm...

$i = 1$; does P_1 agree with **Step 2**? No.

$i = 2$; does P_2 agree with **Step 2**? No.

$i = 3$; does P_3 agree with **Step 2**? Yes. **Work** = **Work**+**Hold**₃; **Finish**₃ = **T**

$i = 1$; does P_1 agree with **Step 2**? Yes. **Work** = **Work**+**Hold**₁; **Finish**₁ = **T**

$i = 2$; does P_2 agree with **Step 2**? Yes. **Work** = **Work**+**Hold**₂; **Finish**₂ = **T**


No more (unfinished) P_i , therefore *safe*.

Example—safe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	2 ³	3 ²	F	2	2 ₁
P ₂	4	1	3	F		
P ₃	2	1	1	F		

Assume P₁ acquires one unit.




Example—safe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	3	2	F	1	1
P ₂	4	1	3	F		0
P ₃	2	1	1	F		0

P₃ can acquire the last unit and finish.




Example—safe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P_1	5	5	0	F	2	\emptyset
P_2	4	1	3	F		
P_3	2	0	0	T		

Then, P_1 can acquire two more units and finish.




Example—safe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	0	0	T	5	3 ₂
P ₂	4	1 ⁴	3 ⁰	F		
P ₃	2	0	0	T		

Finally, P₂ can acquire three more units and finish.



Example—unsafe allocation

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	2	3	F	2	2 ₁
P ₂	5	1	4	F		
P ₃	2	1	1	F		

Assume P₂ acquires one unit.

As before, P₃ can finish and release its resources.

BUT...

i = 1; does P₁ agree with **Step 2**? No.

i = 2; does P₂ agree with **Step 2**? No.

i = 3; does P₃ agree with **Step 2**? Yes. **Work = Work + Hold₂; Finish₂ = T**

Any more unfinished P_i? Yes.


P₁ and P₂ cannot finish. Therefore *unsafe*.

Example—unsafe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	2	3	F	2	2
P ₂	5	2	3	F		
P ₃	2	0	0	T		

NOW...



Resource allocation graph algorithm

- *Banker's Algorithm* works with any number of multiple-instance resources -- but it's a resource intensive algorithm
- If we can assume single instances of each resource, we can use the *Resource Allocation Graph Algorithm*
- This adds a new type of edge to the RAG—a claim edge that indicates that P_i might place a claim on R_j sometime in the future
- Claim edge is actually converted to a request edge when the request is actually made. When a resource is released it is converted back to a claim edge

Resource allocation graph algorithm

- Keeps track of requests coming in the future
 - ◆ all claim edges that will ever occur are present at the start, but we can relax this and add new ones if all edges out of a process are currently only claim edges—that is, the process holds nothing.
- When a process P_i requests a resource, we change the claim edge to a request edge, and then look at what would happen if we turned it into an assignment edge.
- If there's no cycle, we have a safe state—if we have a cycle, it's unsafe and the process must wait.

Deadlock detection

- This technique
 - ◆ does not attempt to prevent deadlocks;
 - ◆ instead, it lets them occur.
- The system
 - ◆ detects when this happens, by periodically running an algorithm to detect a **circular wait** condition,
 - ◆ then takes some action to recover after the fact.
- With deadlock detection, requested resources are
- granted to processes whenever possible.

Deadlock detection

cont.

- A check for deadlock can be made
 - ◆ as frequently as for each resource request, or
 - ◆ less frequently, depending on how likely it is for a deadlock to occur.
- Checking at each resource request has two advantages:
 - ◆ it leads to early detection, and
 - ◆ the algorithm is relatively simple because it is based on incremental changes to the state of the system.
- On the other hand,
 - ◆ such frequent checks consume considerable processor time.

Recovering from deadlocks

- Once the deadlock algorithm has successfully detected a deadlock, some strategy is needed for recovery. There are various ways:
 - ◆ Recovery through *Preemption*

In some cases, it may be possible to temporarily take a resource away from its current owner and give it to another.
 - ◆ Recovery through *Rollback*

If it is known that deadlocks are likely, one can arrange to have processes *checkpointed* periodically. For example, can undo transactions, thus free locks on database records.
 - ◆ Recovery through *Termination*

The most trivial way to break a deadlock is to kill one or more processes. One possibility is to kill a process in the cycle. Warning! *Irrecoverable losses may occur, even if this is the least advanced process.*

Summary of strategies

<i>Principle</i>	<i>Resource Allocation Strategy</i>	<i>Different Schemes</i>	<i>Major Advantages</i>	<i>Major Disadvantages</i>
<i>DETECTION</i>	<ul style="list-style-type: none">• Very liberal; grant resources as requested.	<ul style="list-style-type: none">• Invoke periodically to test for deadlock.	<ul style="list-style-type: none">• Never delays process initiation.• Facilitates on-line handling.	<ul style="list-style-type: none">• Inherent preemption losses.
<i>PREVENTION</i>	<ul style="list-style-type: none">• Conservative; under-commits resources.	<ul style="list-style-type: none">• Requesting all resources at once.• Preemption• Resource ordering	<ul style="list-style-type: none">• Works well for processes with single burst of activity.• No preemption is needed.• Convenient when applied to resources whose state can be saved and restored easily.• Feasible to enforce via compile-time checks.• Needs no run-time computation.	<ul style="list-style-type: none">• Inefficient.• Delays process initiation.• Preempts more often than necessary.• Subject to cyclic restart.• Preempts without immediate use.• Disallows incremental resource requests.
<i>AVOIDANCE</i>	<ul style="list-style-type: none">• Selects midway between that of detection and prevention.	<ul style="list-style-type: none">• Manipulate to find at least one safe path.	<ul style="list-style-type: none">• No preemption necessary.	<ul style="list-style-type: none">• Future resource requirements must be known.• Processes can be blocked for long periods.