


Caching

The Basic Idea




COMP 512: Caching

Why Web Caching?

- ❑ Client
 - ☆ Reduce Response Time
 - fast local access similar to replication
- ❑ Server:
 - ☆ Reduce Load
 - load distribution between cache and server
- ❑ Network:
 - ☆ Save bandwidth

COMP 512: Caching

Where to Cache?



- ❑ Browser
 - ☆ Small: 10MBs memory, 100MB disk
 - Recursive caching (memory vs. disk)
 - 20% hit rate
- ❑ Organization (client-side proxy)
 - ☆ Large: GByte with disk
 - ☆ 50% hit rate
- ❑ Cache-Server Organization
 - ☆ Between Clients and Servers
- ❑ In front of server
 - ☆ Large: GByte
- ❑ Server itself (in memory)

COMP 512: Caching

Why does it work?

- ❑ Requested object in cache because object was previously requested
 - ☆ True *ONLY* for popular objects
- ❑ Works if there are relatively FEW objects that are requested FREQUENTLY
 - ☆ Popular objects are likely to be cached
 - ☆ Not so popular objects are very unlikely to be cached
- ❑ Hit rates that can be reached: 50%
- ❑ Note: miss leads to more messages due to indirection

COMP 512: Caching

When does it not work?

- ❑ Capacity
 - ☆ Not enough cache (object was in cache but purged before rerequested)
 - ☆ Solution: bigger cache, distributed cache....
- ❑ First access to an object
 - ☆ Solution: prefetching
- ❑ Consistency
 - ☆ Object has changed
 - ☆ Cache consistency has its own overhead
- ❑ Dynamic Objects
 - ☆ Dynamic web caching

COMP 512: Caching

Cache Consistency

- ❑ If object at server changes, cached copies become stale
- ❑ Cache consistency mechanisms
 - ☆ Server Push
 - ☆ Client Pull

COMP 512: Caching

Push

- ❑ For each object o of server s , s keeps track of set P of proxies that have requested o
- ❑ When o is modified, notify all proxies in P
 - ☆ Invalidate
 - ☆ Update
 - ☆ Adv. / disadv.?
- ❑ Strong consistency
 - ☆ compare to strong consistency in context of replication

COMP 512: Caching

Pull

- ❑ For each object o cached at proxy p , p polls server s whether o was modified (http: if-modified-since)
- ❑ Periodically: weak consistency
 - ☆ How to determine interval?
- ❑ Whenever object is locally requested: strong consistency (why cache at all in this case?)

COMP 512: Caching

Web-Caching vs. Content Distribution Networks

- ❑ Companies (like Akamai) *replicate* Web sites
 - ☆ Host all (or part) of a Web site for a content provider
 - ☆ Place replicas all over the world on many machines

COMP 512: Caching

Replication / Caching

- ❑ **When:**
 - ☆ replication: planned
 - ☆ cache: on-demand after first request
- ❑ **who updates:**
 - ☆ replication: primary copy vs. update anywhere
 - ☆ cache: server
- ❑ **granularity:**
 - ☆ replication: coarse, e.g., database table, entire database
 - ☆ cache: web-page; fragment; query results; object
- ❑ **origin:**
 - ☆ replication: data of one organization
 - ☆ cache: any data that is accessed
- ❑ **purpose:**
 - ☆ replication: scalability, fast local access fault-tolerance
 - ☆ cache: fast local access; scalability

COMP 512: Caching

Replication / Caching

	Data Replication	Cache
When	planned	on-demand after first request
Who updates	primary copy vs. update anywhere	server
Consistency	weak and strong consistency	weak consistency
Granularity	coarse: table	web-page; fragment; data records
Origin	data of one organization	cache data from any organization
Purpose	scalability; fast local access; fault-tolerance	fast local access; scalability

COMP 512: Caching