


Scanning

COMP 520: Compiler Design (4 credits)

Professor Laurie Hendren

hendren@cs.mcgill.ca



Readings

Crafting a Compiler:

- Chapter 2, A simple compiler
- Chapter 3, Scanning - Theory and Practice

Modern Compiler Implementation in Java:

- Chapter 1, Introduction
- Chapter 2, Lexical Analysis

Flex tool:

- Manual - `http://flex.sourceforge.net/manual/`
- Reference book, Flex & bison -
`http://mcgill.worldcat.org/title/flex-bison/oclc/457179470`

Background (1), from "Crafting a Compiler"





Figure 3.1: Components of a finite automaton drawing and their use to construct an automaton that recognizes $(a b c^+)^+$.

Background (2) , from "Crafting a Compiler"

(b)

State	Character				
	/	Eol	a	b	...
1	2				
2	3				
3	3	4	3	3	3
4					

Figure 3.2: DFA for recognizing a single-line comment. (a) transition diagram; (b) corresponding transition table.

Background (3), from "Crafting a Compiler"

```
/* Assume CurrentChar contains the first character to be scanned */  
State  $\leftarrow$  StartState  
while true do  
    NextState  $\leftarrow$  T[State, CurrentChar]  
    if NextState = error  
    then break  
    State  $\leftarrow$  NextState  
    CurrentChar  $\leftarrow$  READ( )  
if State  $\in$  AcceptingStates  
then /* Return or process the valid token */  
else /* Signal a lexical error */
```

Figure 3.3: Scanner driver interpreting a transition table.

Tokens are defined by *regular expressions*:

- \emptyset , the empty set: a language with no strings
- ϵ , the empty string
- a , where $a \in \Sigma$ and Σ is our alphabet
- $M|N$, alternation: either M or N
- $M \cdot N$, concatenation: M followed by N
- M^* , zero or more occurrences of M

where M and N are both regular expressions.

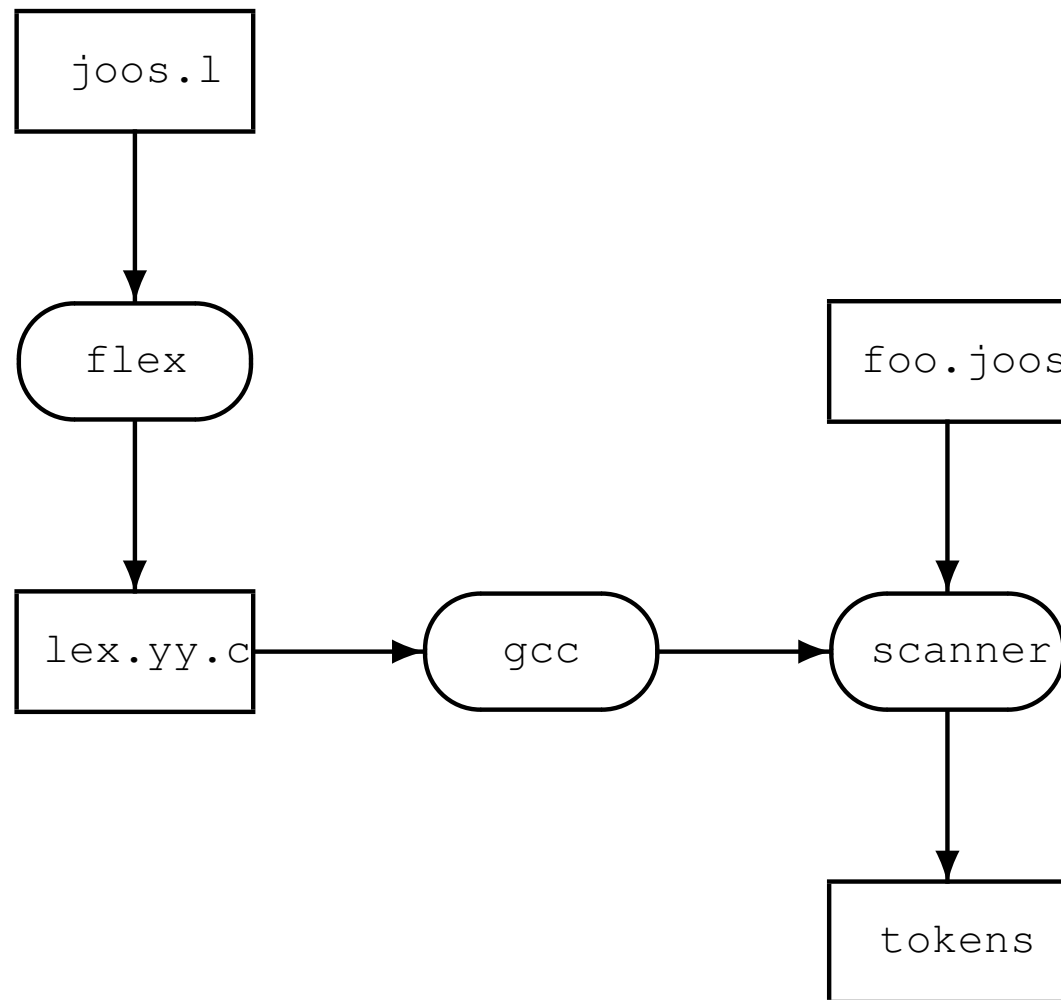
What are M^+ and $M^?$?

We can write regular expressions for the tokens in our source language using standard POSIX notation:

- simple operators: `" * ", " / ", " + ", " - "`
- parentheses: `" (", ") "`
- integer constants: `0 | ([1 - 9] [0 - 9] *)`
- identifiers: `[a - z A - Z _] [a - z A - Z 0 - 9 _] *`
- white space: `[_ \t \n] +`

A *scanner* or *lexer* transforms a string of characters into a string of tokens:

- uses a combination of *deterministic finite automata* (DFA);
- plus some glue code to make it work;
- can be generated by tools like `flex` (or `lex`), `JFlex`, ...



How to go from regular expressions to DFAs?

- `flex` accepts a list of regular expressions (regex);
- converts each regex internally to an NFA (Thompson construction);
- converts each NFA to a DFA (subset construction)
- may minimize DFA

(see “Crafting a Compiler” , ch 3; or “Modern Compiler Implementation in Java”, Ch. 2)

Regular Expressions to NFA (1) from text, "Crafting a Compiler"




Figure 3.17: An NFA with two a transitions.




Figure 3.18: An NFA with a λ transition.

Regular Expressions to NFA (2)from text, "Crafting a Compiler"




Figure 3.19: NFAs for a and λ .




Figure 3.20: An NFA for $A \mid B$.

Regular Expressions to NFA (3)from text, "Crafting a Compiler"




Figure 3.21: An NFA for AB .


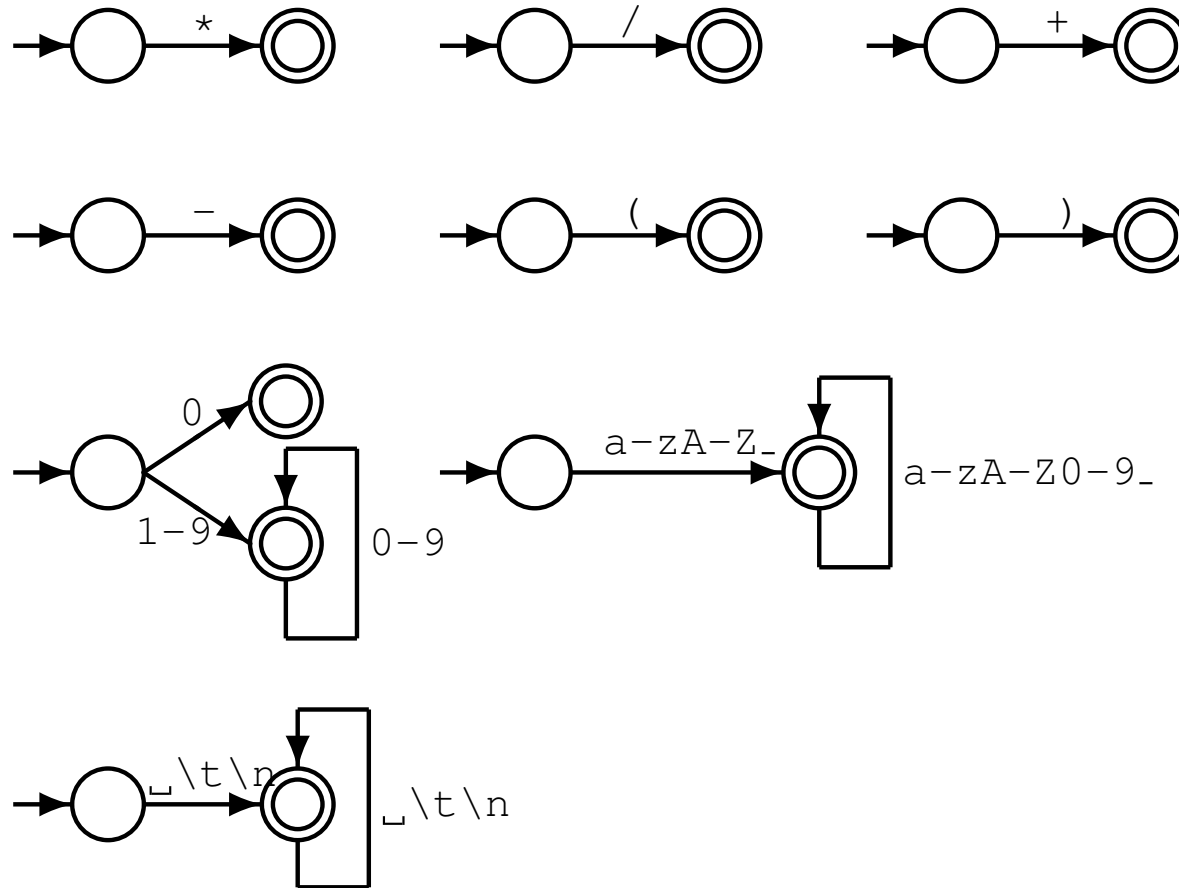


Figure 3.22: An NFA for A^* .



Some DFAs

Each DFA has an associated *action*.

Let's assume we have a collection of DFAs, one for each lex rule

reg_expr1 -> DFA1

reg_expr2 -> DFA2

...

reg_rexpn -> DFA_n

How do we decide which regular expression should match the next characters to be scanned?

Given DFAs D_1, \dots, D_n , ordered by the input rule order, the behaviour of a `flex`-generated scanner on an input string is:

```
while input is not empty do
     $s_i :=$  the longest prefix that  $D_i$  accepts
     $l := \max\{|s_i|\}$ 
    if  $l > 0$  then
         $j := \min\{i : |s_i| = l\}$ 
        remove  $s_j$  from input
        perform the  $j^{\text{th}}$  action
    else (error case)
        move one character from input to output
    end
end
```

- The *longest* initial substring match forms the next token, and it is subject to some action
- The *first* rule to match breaks any ties
- Non-matching characters are echoed back

Why the “longest match” principle?

Example: keywords

```
[ \t]+
    /* ignore */;
...
import
    return tIMPORT;
...
[a-zA-Z_][a-zA-Z0-9_]* {
    yylval.stringconst = (char *)malloc(strlen(yytext)+1);
    printf(yylval.stringconst, "%s", yytext);
    return tIDENTIFIER; }
```

Want to match ```importedFiles``` as `tIDENTIFIER(importedFiles)` and not as `tIMPORT tIDENTIFIER(edFiles)`.

Because we prefer longer matches, we get the right result.

Why the “first match” principle?

Again — Example: keywords

```
[ \t]+
    /* ignore */;
...
continue
    return tCONTINUE;
...
[a-zA-Z_][a-zA-Z0-9_]* {
    yylval.stringconst = (char *)malloc(strlen(yytext)+1);
    printf(yylval.stringconst, "%s", yytext);
    return tIDENTIFIER; }
```

Want to match `‘‘continue foo’’` as `tCONTINUE tIDENTIFIER(foo)` and not as `tIDENTIFIER(continue) tIDENTIFIER(foo)`.

“First match” rule gives us the right answer: When both `tCONTINUE` and `tIDENTIFIER` match, prefer the first.

When “first longest match” (flm) is not enough, look-ahead may help.

FORTRAN allows for the following tokens:

`.EQ., 363, 363., .363`

flm analysis of `363.EQ.363` gives us: `tFLOAT(363) E Q tFLOAT(0.363)`

What we actually want is: `tINTEGER(363) tEQ tINTEGER(363)`

`flex` allows us to use look-ahead, using `' / '`:

`363/.EQ. return tINTEGER;`

Another example taken from FORTRAN, FORTRAN ignores whitespace

1. `DO5I = 1.25` \leadsto `DO5I=1.25`

in C: `do5i = 1.25;`

2. `DO 5 I = 1, 25` \leadsto `DO5I=1, 25`

in C: `for (i=1; i<25; ++i) { ... }`

(5 is interpreted as a line number here)

Case 1: flm analysis correct:

`tID(DO5I) tEQ tREAL(1.25)`

Case 2: want:

`tDO tINT(5) tID(I) tEQ tINT(1) tCOMMA tINT(25)`

Cannot make decision on `tDO` until we see the comma, look-ahead comes to the rescue:

`DO/({letter}|{digit})*=({letter}|{digit})*, return tDO;`

```
$ cat print_tokens.l # flex source code

/* includes and other arbitrary C code */
%{
#include <stdio.h> /* for printf */
%}
/* helper definitions */
DIGIT [0-9]
/* regex + action rules come after the first %% */
%%
[ \t\n]+      printf ("white space, length %i\n", yyleng);
"*"          printf ("times\n");
"/"          printf ("div\n");
"+"          printf ("plus\n");
"-"          printf ("minus\n");
"("          printf ("left parenthesis\n");
")"          printf ("right parenthesis\n");

0|([1-9]{DIGIT}*) printf ("integer constant: %s\n", yytext);
[a-zA-Z_][a-zA-Z0-9_]* printf ("identifier: %s\n", yytext);
%%
/* user code comes after the second %% */
main () {
    yylex ();
}
```

Using `flex` to create a scanner is really simple:

```
$ emacs print_tokens.l
```

```
$ flex print_tokens.l
```

```
$ gcc -o print_tokens lex.yy.c -lfl
```

When input $a * (b - 17) + 5 / c$:

```
$ echo "a*(b-17) + 5/c" | ./print_tokens
```

our print_tokens scanner outputs:

```
identifier: a
times
left parenthesis
identifier: b
minus
integer constant: 17
right parenthesis
white space, length 1
plus
white space, length 1
integer constant: 5
div
identifier: c
white space, length 1
```

Count lines and characters:

```
%{  
int lines = 0, chars = 0;  
%}  
  
%%  
\n      lines++; chars++;  
.      chars++;  
  
%%  
main () {  
    yylex ();  
    printf ("#lines = %i, #chars = %i\n", lines, chars);  
}
```


Remove vowels and increment integers:

```
%{  
#include <stdlib.h> /* for atoi */  
#include <stdio.h> /* for printf */  
%}  
  
%%  
[aeiouy]      /* ignore */  
[0-9]+        printf ("%i", atoi (yytext) + 1);  
  
%%  
main () {  
    yylex ();  
}
```